

# Extending PostgreSQL

EurOmnis 2025

Alex Clay - Suran Systems, Inc.

# About Me

- Alex Clay
- Omnis developer since 2000
- CEO of Suran since 2019
- Used PostgreSQL since 2010
- Thousands of deployments
- Original turn-key installer for on-prem PG installs
- Now all hosted on Suran's managed servers

# About you

# Some Factors at Suran

- Multi-tenant by individual databases
- Multiple tenants per server
- Multi-platform (Omnis, Rails, Native Mobile)
- Driving logic server-side
- 3rd party integrations
- Single source of truth

# PostgreSQL Glue

- dblink
- Foreign Data Wrappers
- Extensions for HTTP calls
- Bulk loading (COPY)
- Shelling out (run programs on the server)
- Bonus :)

# Extensions

# PostgreSQL Extensions

- Extend PostgreSQL with new functionality
- Many built-in extensions: <https://www.postgresql.org/docs/current/contrib.html>
- Install additional extensions:
  - OS package manager
  - Compile
  - PGXN
- Some require loading shared libraries
- Special consideration for managed cloud platforms, like RDS

# Creating an extension

```
CREATE EXTENSION IF NOT EXISTS [extension] SCHEMA [schema]
CASCADE;
```

```
euromnis_demo=# create extension unaccent;
CREATE EXTENSION
```

```
euromnis_demo=# create schema unaccent;
CREATE SCHEMA
euromnis_demo=# alter extension unaccent set schema unaccent;
ALTER EXTENSION
euromnis_demo=# drop extension unaccent;
DROP EXTENSION
```

# Today's Examples

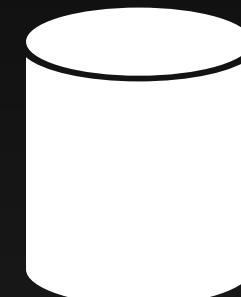
# Example Sources

- <https://github.com/barkingfoodog/euromnis2025-courses>
- Docker
  - Docker Desktop
  - Orbstack
  - colima

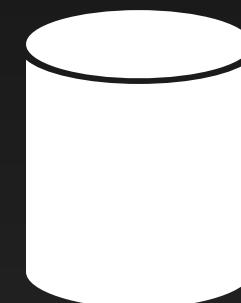


# Use Cases

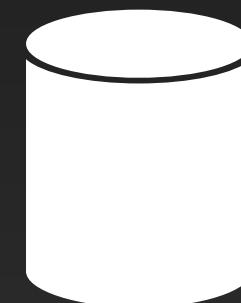
# Use Case: suran\_database\_manager



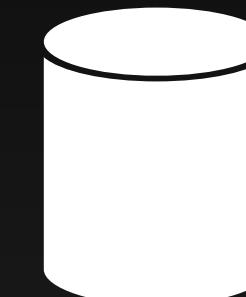
**Client DB 1**



**Client DB 2**



**Client DB 3**



**suran\_database\_manager**

Table of client databases

- Description
- Version
- User ID

How to keep data in suran\_database\_manager up-to-date?

# Use Case: suran\_database\_manager



# dblink

- Used to connect to another PostgreSQL database
- Built into PostgreSQL
- Can run one-off commands
- Can establish a persistent connection
- Best for:
  - Functions
  - One-off SQL commands
- Predecessor to PostgreSQL Foreign Database Wrapper (FDW)

# dblink Connection Types

Type	Advantage	Disadvantage
<b>One-off</b>	+ Least code	- Only runs one SQL command then disconnects
<b>Unnamed</b>	+ Less code + Persistent	- Not explicit - Could cause bugs if using multiple unnamed connections
<b>Named</b>	+ Persistent + Can have multiple, concurrent	- A little more setup - Need to manage your connection

# dblink One-off

```
# select * from dblink.dblink(
  'dbname=geodata',
  'SELECT geo.country_distance(''US'', ''NL'')'
) as t(
  distance float
);
```

distance

---

7514531.024600981  
(1 row)

# dblink Unnamed

```
euromnis_demo=# select dblink.dblink_connect('dbname=geodata');
dblink_connect
-----
0K
(1 row)

euromnis_demo=# select * from dblink.dblink('SELECT geo.country_distance(''US'', ''NL'')') as t(distance float);
distance
-----
7514531.024600981
(1 row)

euromnis_demo=# select dblink.dblink_disconnect();
dblink_disconnect
-----
0K
(1 row)
```

# dblink Named

```
euromnis_demo=# select dblink.dblink_connect('geo', 'dbname=geodata');
dblink_connect
-----
0K
(1 row)

euromnis_demo=# select * from dblink.dblink('geo', 'SELECT geo.country_distance(''US'', ''NL'')')
as t(distance float);
distance
-----
7514531.024600981
(1 row)

euromnis_demo=# select dblink.dblink_disconnect('geo');
dblink_disconnect
-----
0K
(1 row)
```

# Using dblink

- dblink() to query and return results
  - MUST specify return structure
- dblink\_exec() to send a command without results
  - void functions
  - INSERTS
- dblink\_send\_query() sends asynchronous connection
- Cursor functions are available as well

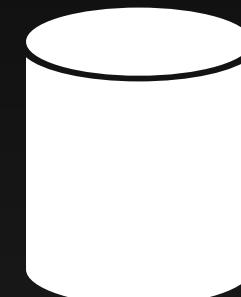
# Using dblink

- Connection uses libpq, same as Omnis PG DAM
- Can pass configuration values, sslmode, etc.
- Authentication via .pgpass or foreign server/user mapping (more on that later)
- <https://www.postgresql.org/docs/current/dblink.html>
- <https://www.postgresql.org/docs/current/libpq-connect.html#LIBPQ-CONNSTRING>

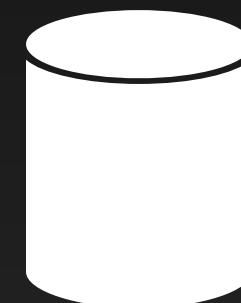
# Demo

```
geo.country_distance(source_country, destination_country)  
database_utilities.update_fgdb_info()
```

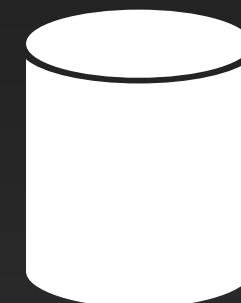
# Use Case: Processor Data Exchange



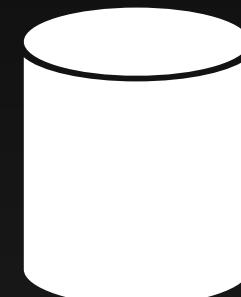
**Client DB 1**



**Client DB 2**



**Client DB 3**



**pde**

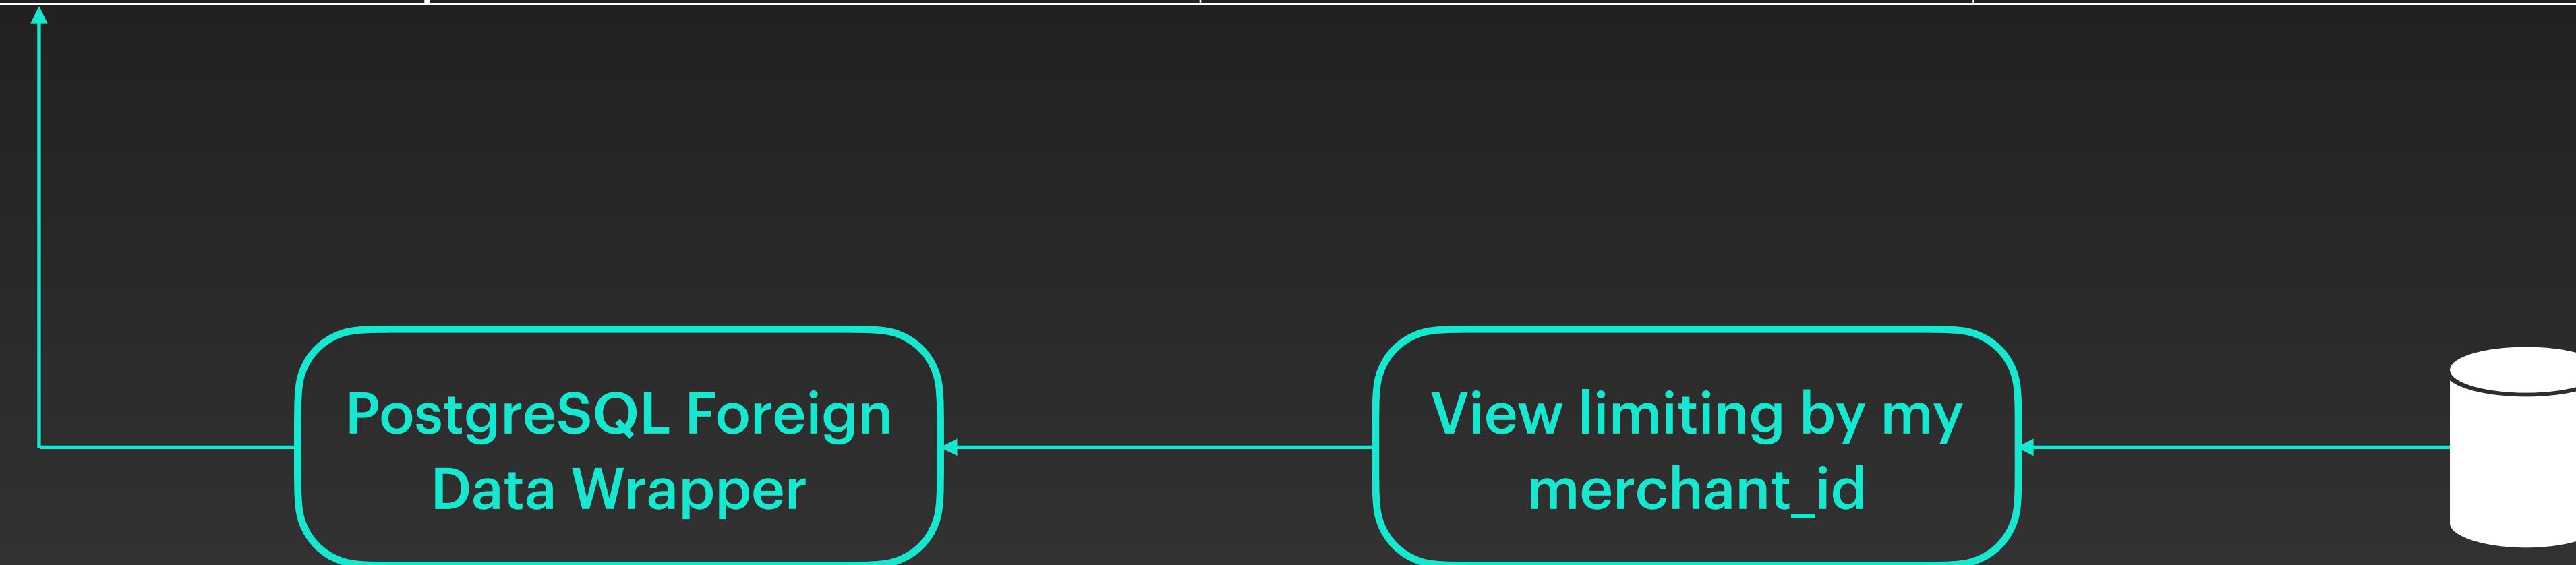
Transaction data for all clients

Organized by merchant\_key

How to access filtered data and run repeated queries?

# Use Case: Processor Data Exchange

merchant_key	payment_id	date	amount
16	523	2025-10-01	\$102.34
16	578	2025-10-02	\$56.98
23	580	2025-10-02	\$230.21



# Foreign Data Wrapper (FDW)

- SQL/MED standard
- Exposes an external or foreign data source via SQL
- Use standard CRUD (SELECT/INSERT/UPDATE/DELETE)
- Built-in:
  - `postgres_fdw` (Access other PostgreSQL or even your own!)
  - File (Query a file or program, alternative to COPY)
- MANY other options!
- [https://wiki.postgresql.org/wiki/Foreign\\_data\\_wrappers](https://wiki.postgresql.org/wiki/Foreign_data_wrappers)

# FDW Components

## EXTENSION

Which FDW?

## SERVER

How to get to the data  
—host, database, file  
path, URL, etc.

## USER MAPPING

Maps local PG roles  
(users) to credentials  
for the server

## FOREIGN TABLE/SCHEMA

Local relation (table) representing  
the remote data; Can import a full  
schema for convenience

# PostgreSQL FDW

- Maintained with the core PostgreSQL code
- Improvements with each release, especially for push-down
- Supports push-down where possible
  - Normally, queries/joins pull the full data set from the foreign server, then join or reduce results locally
  - If possible, `postgres_fdw` asks the remote server to query/join (plan and execute) then only transfers necessary data
    - Big performance benefit
    - Some SQL is too complex for this

# Materialized Views

- GREAT way to manage FDW performance, especially over a remote connection for read-only operations
- Strategy:
  - Set up foreign table
  - Create a materialized view for that table
  - Query/join work with materialized view
  - Refresh if/when needed

# Materialized Views

```
-- Foreign Table pde.tempe_transaction_auth  
# select * from pde.tempe_transaction_auth where  
merchant_key = '33551';  
Time: 806.358 ms
```

```
-- Materialized View pde.transaction_auth  
# select * from pde.transaction_auth where merchant_key =  
'33551';  
Time: 179.144 ms
```

# PostgreSQL FDW

- Transactions are created on the remote server to mirror the local connection
- This solves a major challenge with distributed data updates
- Connection management is normally pretty automatic
  - Connect when needed, once per session/server/user
  - Closed when your session terminates
- Can manage with `keep_connections` or specific functions (see docs)

# Compare with dblink

- No function calls with FDW unless mapped to a view
- FDW can be MUCH more performant
- Easier to join and work with foreign tables
- Importing a foreign schema avoids needing to maintain table definition in two places (though PG 18 adds a LIKE option)
- FDW has more setup than dblink, but it persistent and self-healing
- Transaction management is much better with FDW

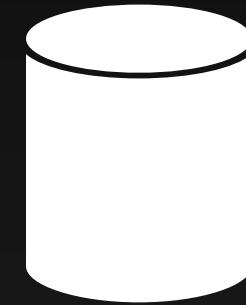
# Demo

geo.country

pde

ucc-data-hub-import

# Use Case: API integration from PG



**Client DB**



**Suran API**

Running on same host as PG

How to call API endpoints from database, e.g. triggers?

# Use Case: API integration from PG



**Client DB**



**Atlassian Jira**

Stores courtesy contact issues

Used by Suran to follow up  
with clients we haven't  
engaged with recently

How to automatically create new courtesy contacts?

# HTTP from PostgreSQL

- `pgsql-http` extension
- Wraps `libcurl` for HTTP calls
- Basic `http()` function takes a request and returns a response
- Various convenience wrapper functions with methods and content types
- Handy `urlencode()`
- Talk to any web service, including Omnis!

# Example HTTP Call

```
euromnis_demo=# select * from http.http_get('https://api.ipify.org?  
format=json');  
-[ RECORD 1 ]  
+-----  
|-----  
|-----  
  
status | 200  
content_type | application/json  
headers | {"(date,\\"Sun, 05 Oct 2025 07:28:15 GMT\\")","(content-type,application/json)","(server,cloudflare)","(vary,Origin)","(cf-cache-status,DYNAMIC)","(content-encoding,br)","(cf-ray,989b3381199d8cb0-AMS)"}  
content | {"ip":"83.167.209.169"}
```

# Example HTTP Call

```
# select trim(both '\"' from (content::jsonb -> 'ip')::text) as public_ip from http.http_get('https://api.ipify.org?format=json');  
public_ip
```

---

```
83.167.209.169  
(1 row)
```

# Demo

IP Address

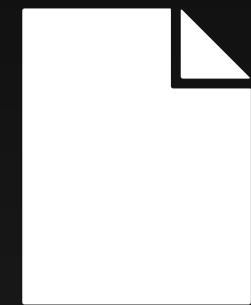
Currency exchange

populate\_courtesy\_contacts

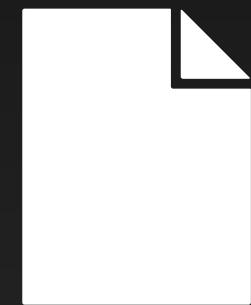
# Use Case: Bulk Loading Data



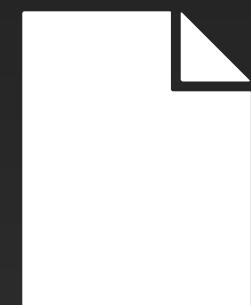
**Client DB**



**CSV Data**



**CSV Data**



**CSV Data**

How to easily and quickly load data into PostgreSQL?

# COPY

- Use COPY to move large amounts of data into and out of PostgreSQL
- Default is a plain text format
- Can also read/write CSV, binary or PROGRAM! (More to come)
- Options for handling NULLs, quoting, error handling
- Can target data to load in out (COPY TO query; add WHERE to COPY FROM)
- Expose as table using file\_fdw

# COPY vs. \copy

Type	Example	Use
<b>COPY</b>	<b>COPY geo.country TO country.csv</b>	<ul style="list-style-type: none"><li>• Works with a server-side file</li><li>• Must be R/W by the postgres process</li></ul>
<b>psql</b>	<b>\copy geo.country TO country.csv</b>	<ul style="list-style-type: none"><li>• Accesses files by psql client</li><li>• More versatile for permissions, remote access</li></ul>

# Timings

Method	Time 1	Time 2	Time 3	Average
<b>COPY (2m rows)</b>	1443ms	1441ms	1464ms	1449ms
<b>\copy (2m rows)</b>	1494ms	1460ms	1466ms	1473ms
<b>Omnis (1k rows)</b>	5340ms	6900ms	7080ms	6440ms

# file\_fdw

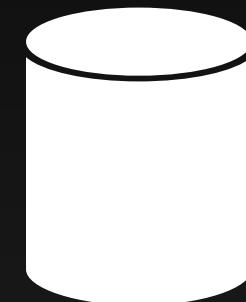
- Expose any file readable by COPY as a SQL table
- Useful for querying without loading into a table
- No indexes or optimizations; use a table if you need better performance
- Define a single server; specify the file on the FOREIGN TABLE

# Demo

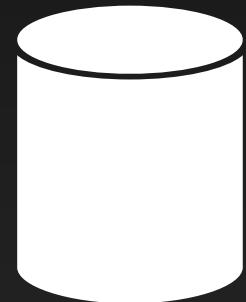
people data

doc import people\_and\_roles

# Use Case: Distributing Schemas



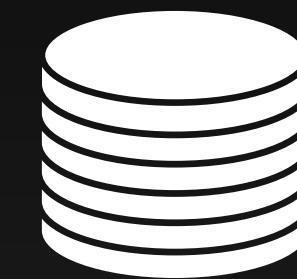
**Client DB 1**



**Client DB 2**



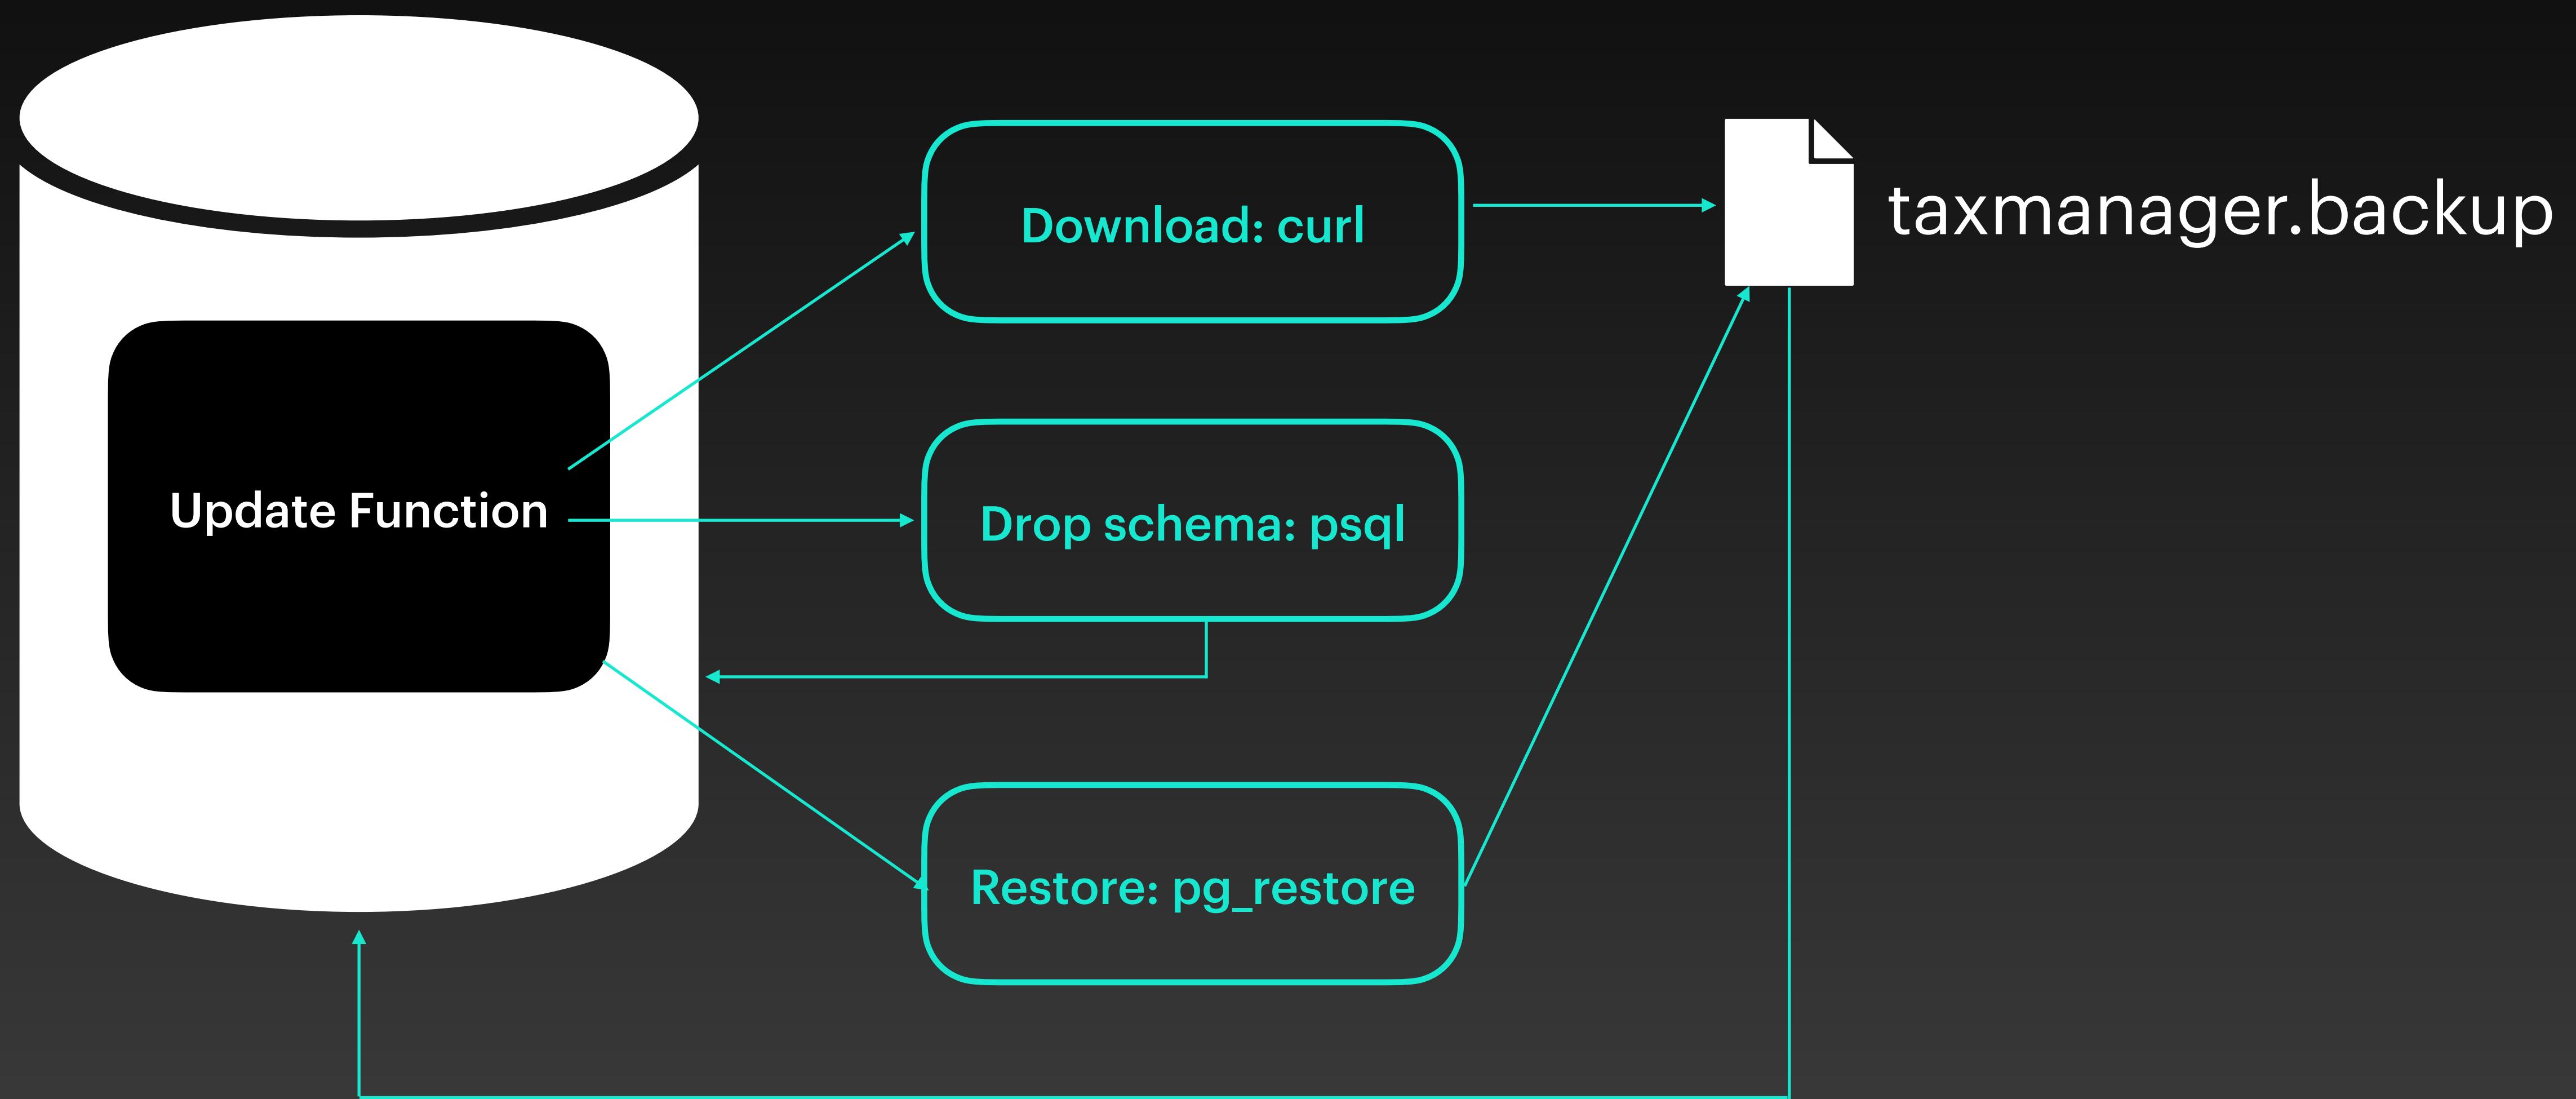
**Client DB 3**



**Static data and functions stored  
in a single schema**

How to ship updated data and functions to all users on demand?

# Use Case: Distributing Schemas



# COPY FROM/TO PROGRAM

- Launches a program server-side as the user running postgres
- Consider working directory, permissions of that user
- Returns stdout for COPY FROM
- Sends stdin for COPY TO
- Can add stderr and return code with a little manipulation
- Can expose as a foreign table using file\_fdw

# Example COPY FROM PROGRAM

```
euromnis_demo=# CREATE TEMPORARY TABLE shell_output (output text);
CREATE TABLE
euromnis_demo=# COPY shell_output FROM PROGRAM 'whoami';
COPY 1
euromnis_demo=# \TABLE shell_output;
      output
-----
postgres
(1 row)

euromnis_demo=# \DROP TABLE shell_output;
\DROP TABLE
```

# Example COPY FROM PROGRAM file\_fdw variant

```
euromnis_demo=# CREATE FOREIGN TABLE shell_output (output text)
 SERVER file_fdw OPTIONS (program 'whoami');
CREATE FOREIGN TABLE
euromnis_demo=# TABLE shell_output;
 output
-----
postgres
(1 row)
```

# Example COPY TO PROGRAM

```
euromnis_demo=# copy (table geo.country) to program 'gzip >
country.csv.gz' with
csv header;
COPY 245
```

```
ea5f13a40fa7:/# gunzip $PGDATA/country.csv.gz
ea5f13a40fa7:/# head -n 3 $PGDATA/country.csv
country,latitude,longitude,name
AD,42.546245,1.601554,Andorra
AE,23.424076,53.847818,United Arab Emirates
```

# Demo

run\_shell\_command

update\_taxmanager

# Bonus

# Fun topics

- pg\_curl: [https://github.com/RekGRpth/pg\\_curl](https://github.com/RekGRpth/pg_curl)
- pg\_background: [https://github.com/vibhorkum/pg\\_background](https://github.com/vibhorkum/pg_background)
- Logical Replication: <https://www.postgresql.org/docs/current/logical-replication-subscription.html#LOGICAL-REPLICATION-SUBSCRIPTION-EXAMPLES>
- pgAgent
- plperlu... 